

CPSC 317

INTERNET COMPUTING

Module 6: Reliable Transmission and TCP

1

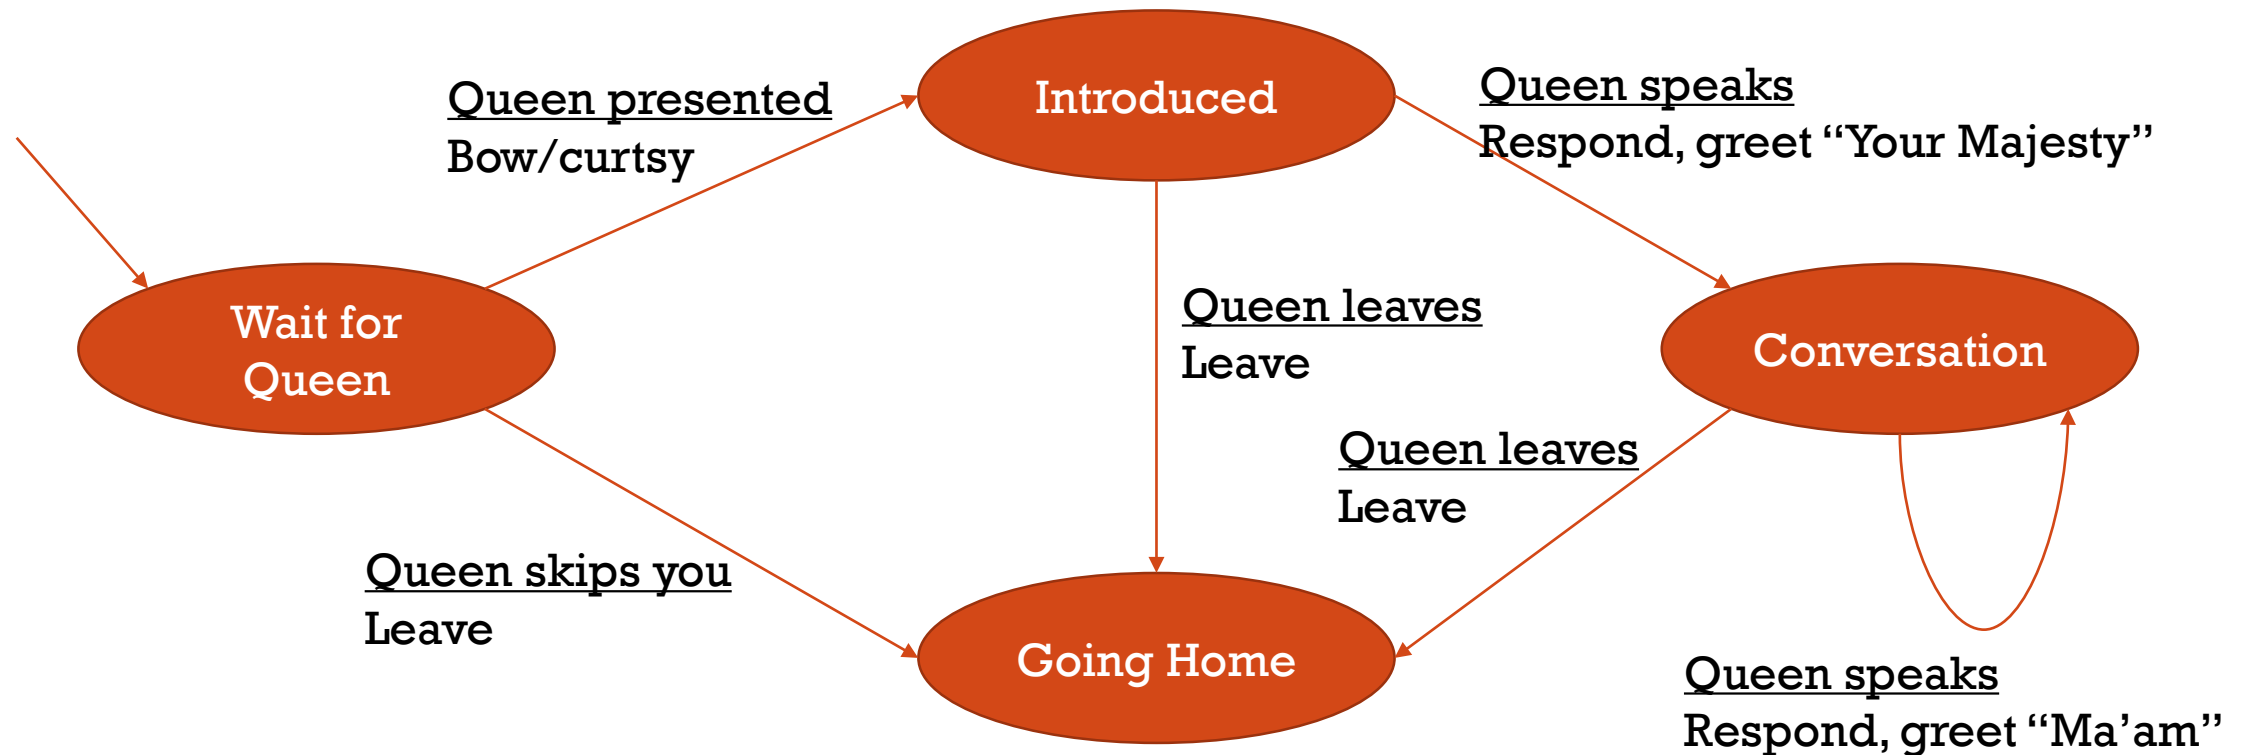
REMINDER: PROTOCOLS

- From textbook: A **protocol** defines the format and the order of messages exchanged between two or more communicating entities, as well as the actions taken on the transmission and/or receipt of a message or other event
- Key concepts:
 - Format and order of messages
 - Actions to be taken on events
- A fully-defined protocol must provide a proper action for any event in any state

PROTOCOL EXAMPLE: GREETING THE QUEEN

- Don't speak until spoken to
- Upon presentation to the Queen
 - Men do a neck bow
 - Women a slight curtsy
- Queen speaks to you
 - First salutation is “Your Majesty”
 - Subsequently it is “Ma'am”

PROTOCOLS AS STATE MACHINES



BUILDING A RELIABLE PROTOCOL

- Let's create a protocol for reliable delivery
 - Send only one packet at a time
 - Identify when sending is allowable action
 - Identify when resending is required
 - Enumerate events and actions for both sender and receiver
 - Draw state machine
- Initial scenario assumptions
 - One sender, one receiver
 - Data to send comes from upper layer
 - Packets are never lost, but may be corrupted

POSSIBLE EVENTS AND ACTIONS

Receiver

- Packet received without problems
 - Send ACK
- Packet received corrupted
 - Send NAK

Sender

- Data ready to send
 - Send data
- Receive ACK
 - Get ready to send more data
- Receive NAK
 - Resend data

STATE MACHINE: CORRUPTION SCENARIO

Receiver

Received valid data

Forward data

Send ACK




Received corrupt data

Send NAK

Sender

Data to send

Send data



Received ACK

^

Received NAK

Resend data

CORRUPT RESPONSE

- What if the ACK/NAK is corrupted?
- What if:
 - We ignore the corrupt ACK/NAK?
 - We treat it as an ACK?
 - We treat it as a NAK?

DEALING WITH DUPLICATE TRANSFER

- How can the receiver deal with a duplicate packet?
- What information needs to be included to allow receiver to identify packet as duplicate?
- Reminder: scenario does not lose data, only corrupts

POSSIBLE EVENTS

Receiver

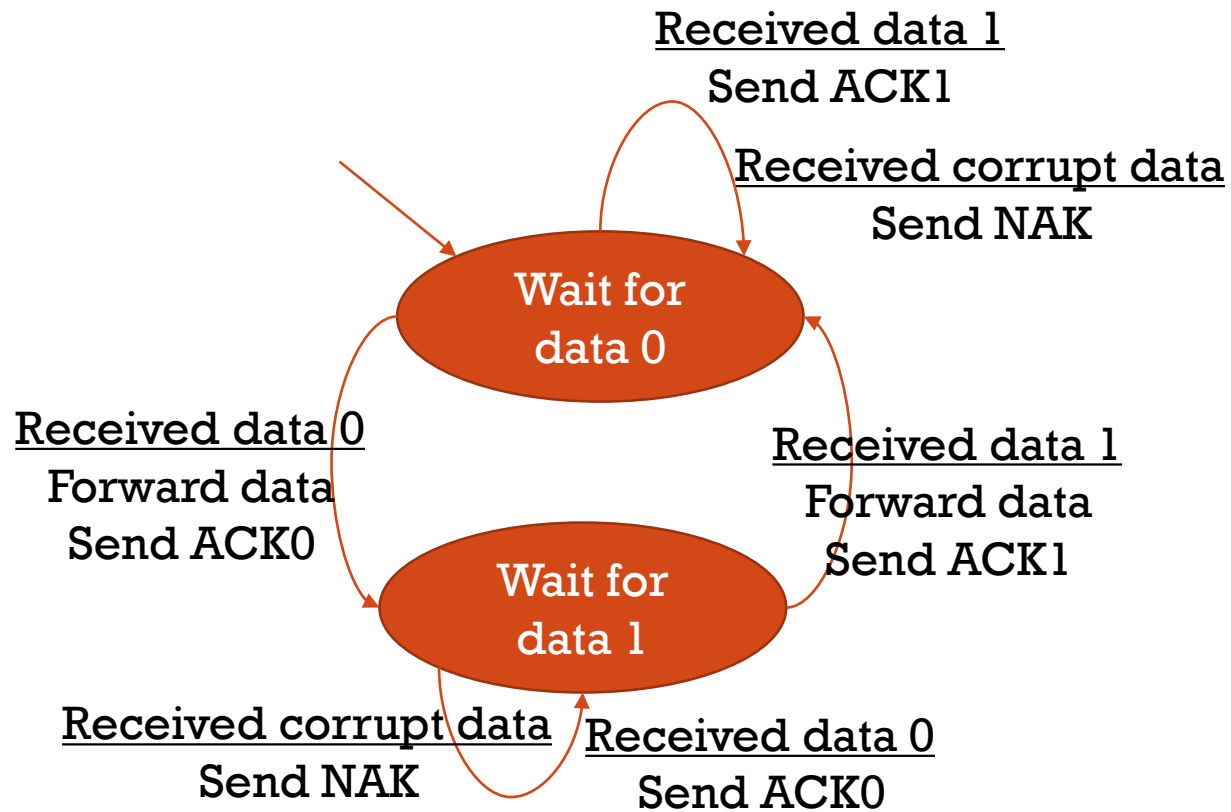
- Packet 0 received without problems
- Packet 0 received corrupted
- Packet 1 received without problems
- Packet 1 received corrupted

Sender

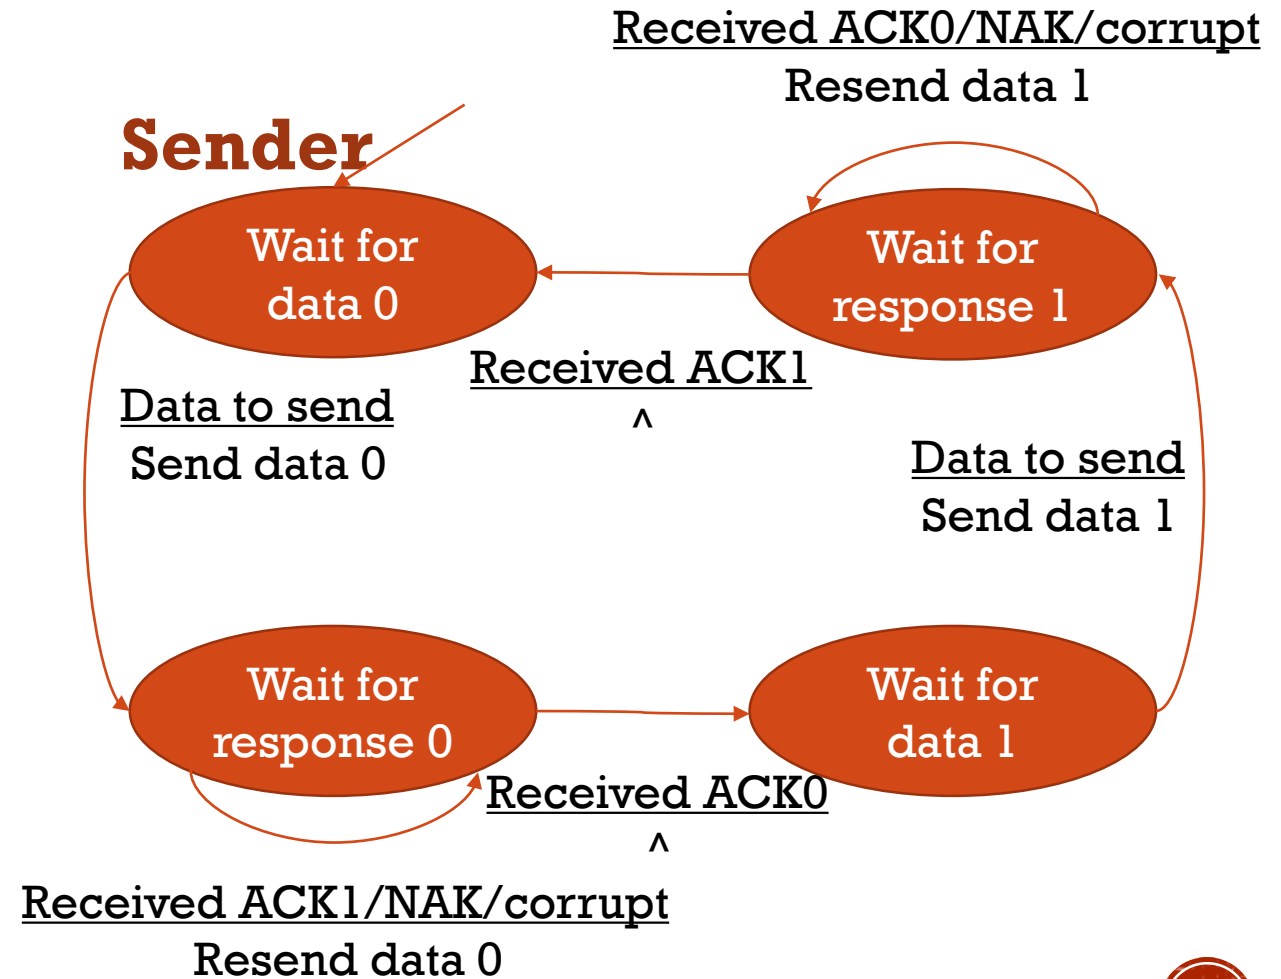
- Data ready to send
- Receive ACK0/ACK1
- Receive NAK
- Receive corrupt response

ALTERNATE BIT PROTOCOL

Receiver

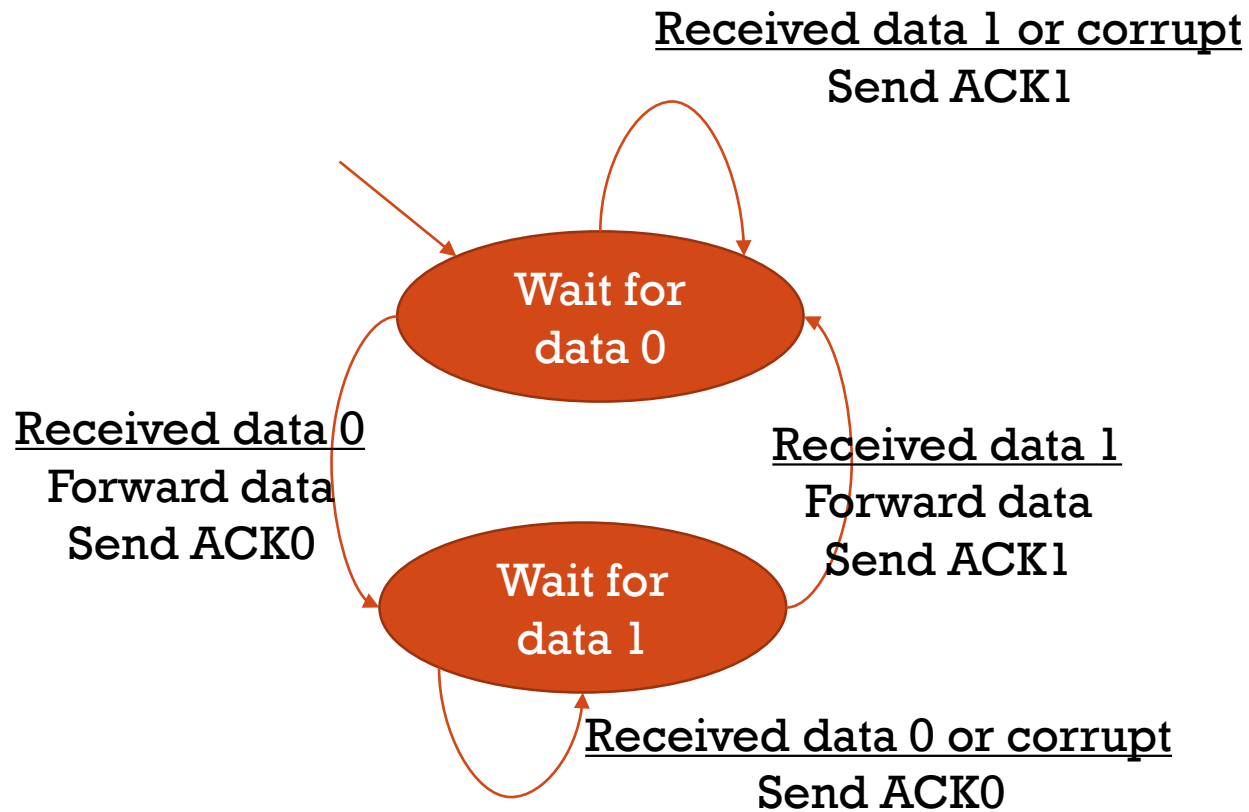


Sender

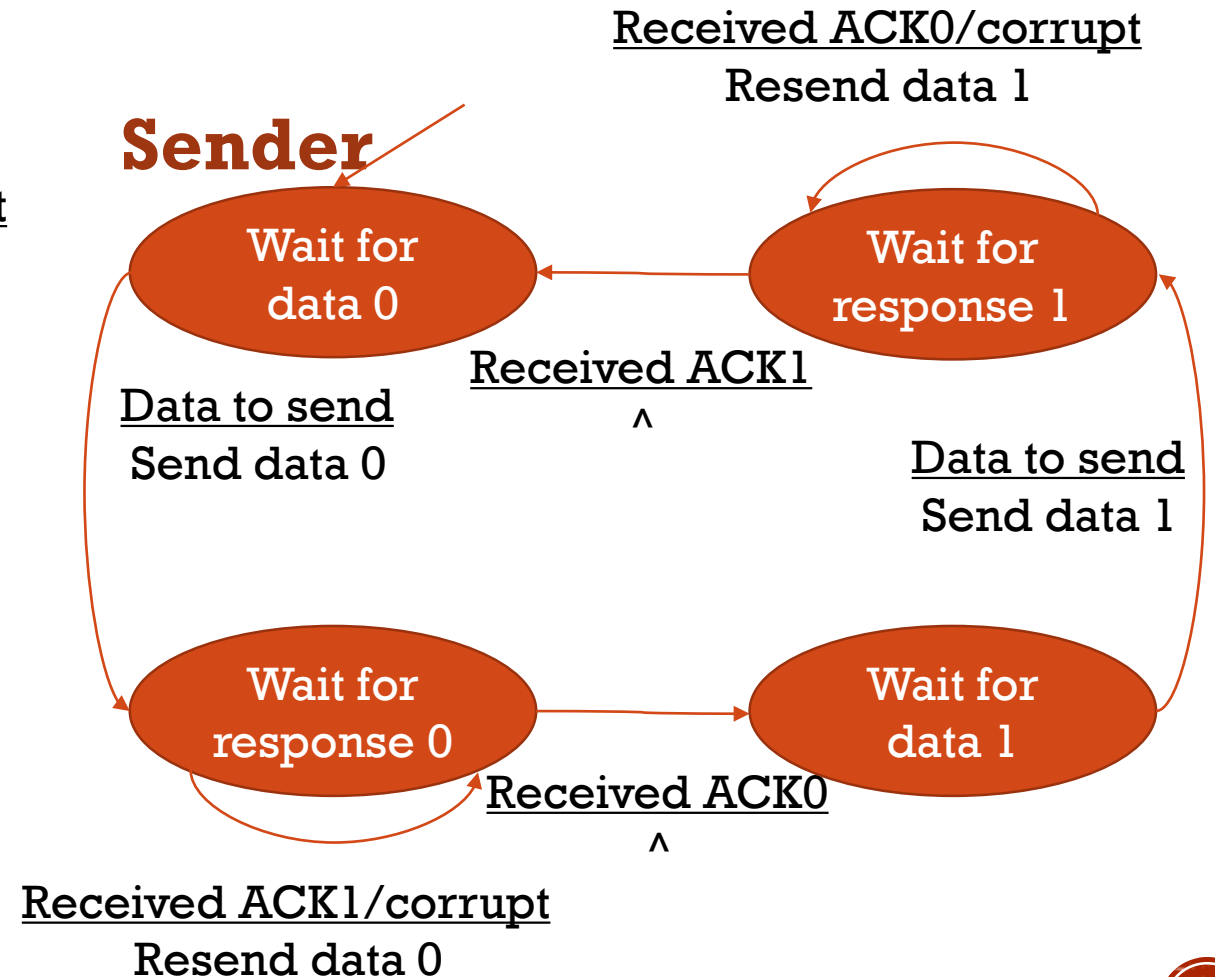


ALTERNATE BIT PROTOCOL (WITHOUT NAK)

Receiver



Sender



WHAT IF A PACKET IS LOST?

- What happens if a packet is lost?
- What happens if an ACK is lost?
- How to determine if a packet is lost?
 - Who determines it? Sender or receiver?

WHAT IF A PACKET IS LOST?

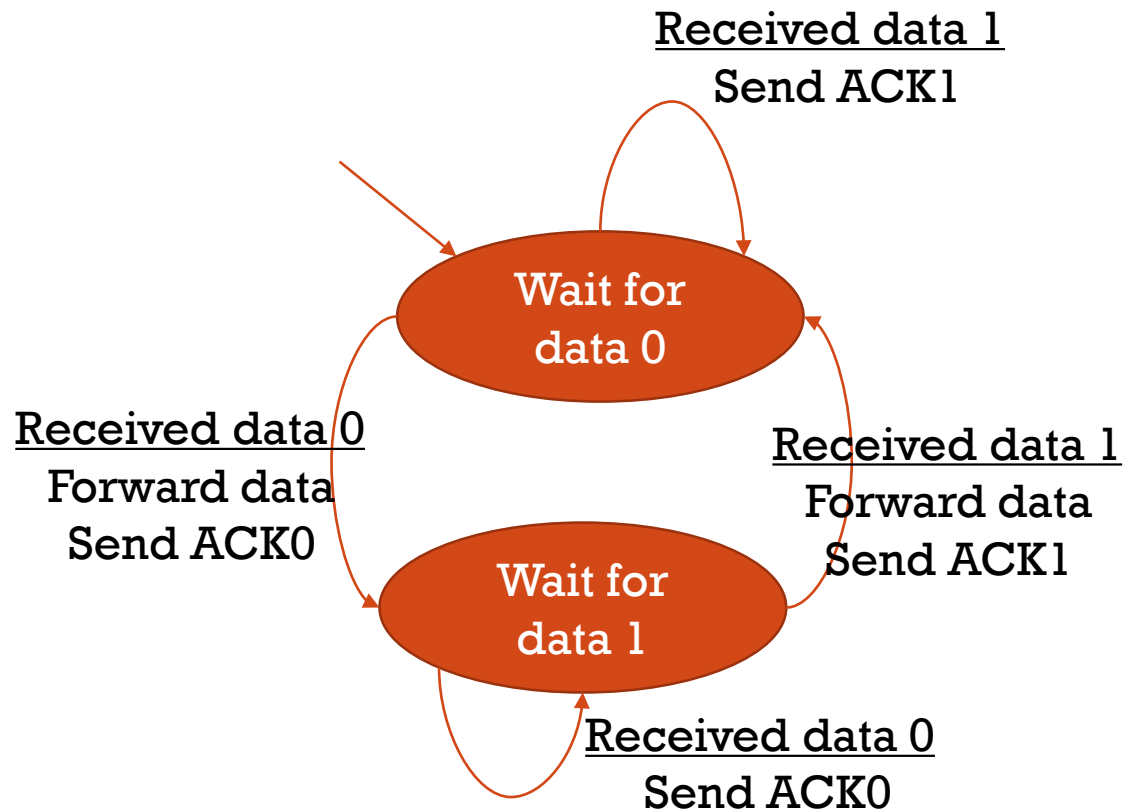
- What happens if a packet is lost?
- What happens if an ACK is lost?
- What changes are needed on the receiver?
- What changes are needed on the sender?

CORRUPT OR LOST?

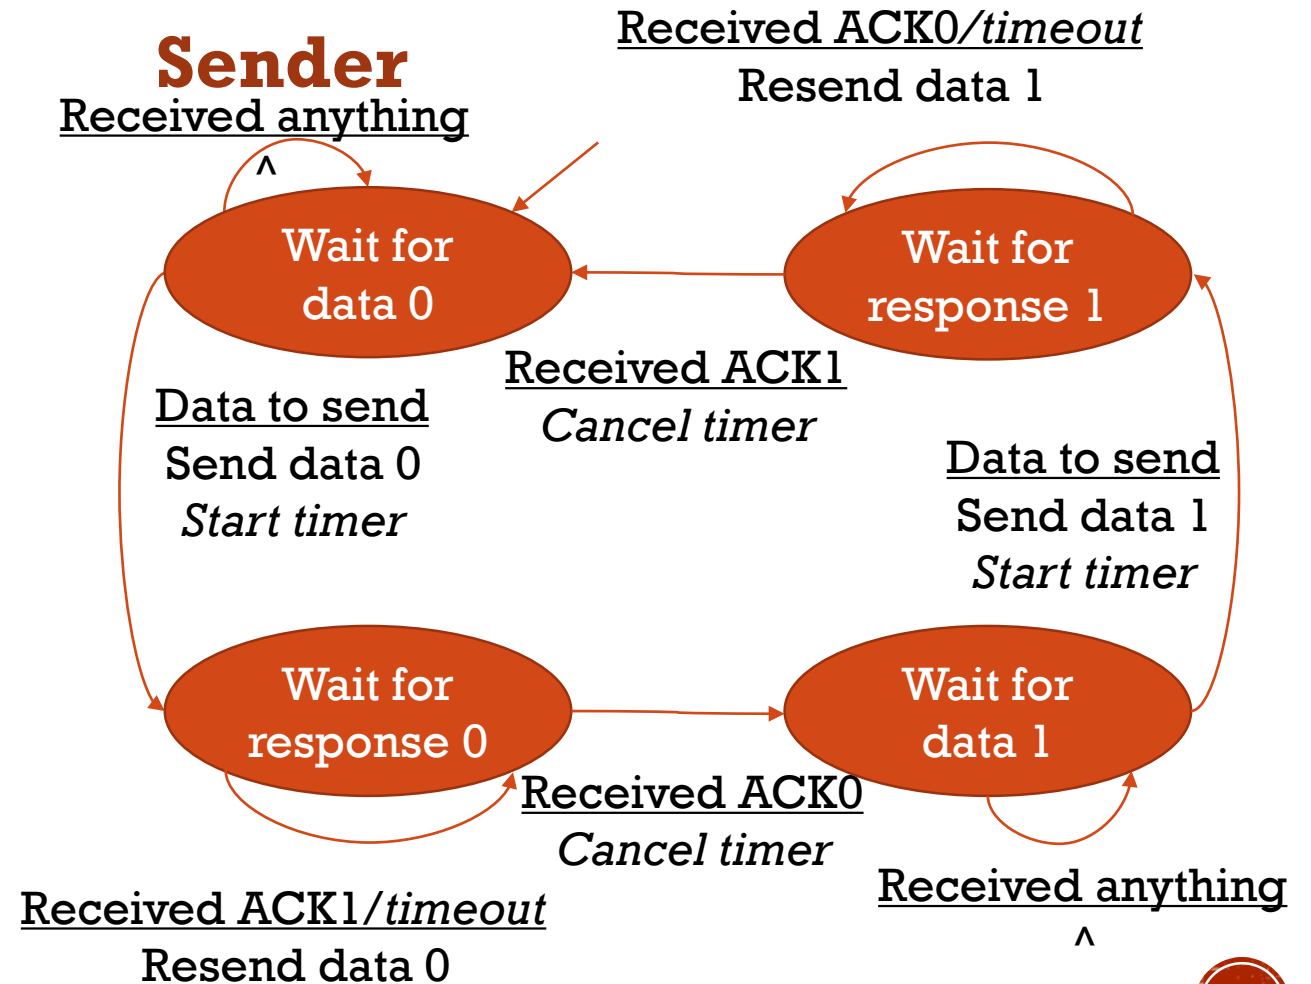
- Does it make sense to distinguish corrupt data from lost data?
 - If something is corrupt, we don't know what's corrupt
 - Headers may have been corrupted too
 - The corrupt packet may not even be for you!!!
 - So, corrupt data becomes lost data (usually at link layer)

ALTERNATE BIT PROTOCOL (TIMEOUT)

Receiver



Sender



HOW LONG SHOULD TIMEOUT BE?

- Should it be the same across connections?
- Should it always be the same for the same connection?
- What happens if timeout is too long?
- What happens if timeout is too short?
- What measured metric can be used to infer a proper timeout?

HOW LONG SHOULD TIMEOUT BE?

- Assume sequential packets take the following sequence of RTT to be ACKed. What should be the timeout?
 - 10, 10, 10, 10, 10, 10, 10, 10
 - 100, 10, 100, 100, 100, 100, 100
 - 1000, 10, 10, 10, 10, 10, 10, 10
 - 10, 100, 10, 100, 10, 100, 10, 100
 - 1, 1, 1, 5, 5, 5, 10, 10, 40, 40, 40, 60, 100, 100, 2, 2, 2, 4, 2

ESTIMATING TIMEOUT

- Simple average of RTTs doesn't respond quickly enough
 - Also doesn't capture jitter
- Timeout value must adapt:
 - Track changes in RTT over time
 - Accommodate packet-to-packet deviations due to jitter
- Discussion in more detail: textbook 3.5.3

TIMEOUT FORMULA

- Assuming measured RTT of t

- Estimated RTT:

$$ERTT_i = (1 - \alpha) \times ERTT_{i-1} + \alpha \times t$$

- Deviation of RTT (captures jitter):

$$\Delta RTT_i = (1 - \beta) \times \Delta RTT_{i-1} + \beta \times |t - ERTT_{i-1}|$$

- Timeout:

$$ERTT_i + 4 \times \Delta RTT_i$$

- Suggested values: $\alpha = 0.125, \beta = 0.25$

ALTERNATE BIT PROTOCOL IN PRACTICE

- Assume a connection from Vancouver to Montreal
 - RTT is 30 ms
 - Link speed is 1Gbps
 - Packet size is 1000 bytes, including overhead
- How much of the bandwidth are we actually using?
 - Only one packet is sent at a time
 - Transmission delay, one packet: $8000/10^9 = 0.008 \text{ ms}$
 - Utilization: $0.008/30+0.008 = 0.00027$ (or 0.027%)
- How can we solve this?

SENDING MULTIPLE PACKETS

- Sender can send multiple packets
- Don't wait for each acknowledgement
 - Assume most packets are successful
- Sender needs to save sent packets to potentially resend
 - Size is limited

SENDER'S WINDOW

- Sender's window: range of packets that are stored for potential resend
- Window only moves when first packet in window is acknowledged
 - What is other ACKs are received?
- New packets are sent only when they “fit” in the window

RECEIVER'S WINDOW

- What if receiver receives packets out of order?
- Receiver's window: store packets received out of order
- Once missing packets arrive, window is processed
- Packets received beyond the window limit will be discarded

PROBLEMS TO CONSIDER

- Sender

- How does the sender know that data got lost?
- Can lost data be distinguished from a lost ACK?
- If we send more than one packet, how many packets can we remember?

- Receiver

- How can you tell if data is out-of-order or missing?
- What should be ACKed?

GO-BACK-N STRATEGY

- Receiver:
 - When packet is received, send ACK for last packet received in order
 - Discard arriving packet if out of order (no receiver window)
- Sender:
 - Can have a specific number of outstanding (unacknowledged) packets in memory: ***sender's window***
 - Start timer on first packet sent
 - On timeout go to last unack'ed packet and resend everything (restart timer)
 - Received ACKs may be cumulative (restart timer on receipt)

SELECTIVE-REPEAT STRATEGY

- Receiver:
 - Each packet is ack'ed individually
 - Out of order packet is stored for later: ***receiver's window***
- Sender:
 - Can have a specific number of outstanding (unacknowledged) packets in memory: ***sender's window***
 - Each packet has its own timer
 - Each packet is individually resent if timeout is reached
 - ACKs received in order move the sender's window

SEQUENCE NUMBER RANGE

- The range of possible sequence numbers is limited by the number of bits used for it
 - Example: 3 bits gets numbers 0-7, 8 bits gets numbers 0-255
 - Sequences return to zero (e.g., in range 0-7, after 7 comes another 0)
- What is the maximum sender window size for the range 0-255?
 - Maybe easier to compute: what about range 0-3?
 - Can the receiver distinguish a new 0 from a resent old 0?
 - Does the answer change for selective repeat vs go-back-N?

SEQUENCE NUMBER RANGE

- Rule: sender's window size + receiver's window size \leq sequence number range
- For a range with n numbers (0 to $n - 1$):
 - Go-back-N:
 - receiver's window size is 1
 - sender's maximum window size is $n - 1$
 - Selective Repeat:
 - any values for window sizes that add up to n is fine
 - for same size on both sides, use $\lfloor n/2 \rfloor$
- Why?

FLOW CONTROL

- Should we always send the full window size?
- What if the receiving application is slow accepting new data?
 - Packets will accumulate in the receiver's buffer
 - Eventually buffer will be full, packets will be dropped
 - Immediately resending this data does not resolve the problem
- Receiver will notify the sender how much data it can handle
 - This information is usually included in the ACK
 - Sender adjusts its window size based on this information

CONGESTION CONTROL

- What if the network can't handle a full window's worth of data?
 - Packets and ACKs will be dropped by the routers
- Missing ACKs are a sign that there is congestion somewhere (in either direction)
- Sender can reduce sending window once congestion is detected
 - Example: if some number of ACKs are missing in a period of time
- If all packets are ACKed, we can increase the window again

EFFECT ON THROUGHPUT

- Throughput is affected by
 - Bandwidth of sender's direct network connection
 - Receiver's specified window size (flow control)
 - Sender's adjusted window size (congestion control)
- If sender's direct connection is not the bottleneck:
 - Another router will experience congestion elsewhere
 - Congestion control will reduce transmission speed

TCP IMPLEMENTATION

- Both sides act as sender and receiver
 - All packets have both a sequence number and an ACK
 - Sequence numbers in one direction correspond to ACK numbers in opposite direction
 - Sequence numbers are incremented by payload size
- Retransmission strategy:
 - ACKs correspond to first sequence number not yet received (similar to Go-Back-N)
 - Receiver stores packets in its own window (like Selective-Repeat)
 - Three or more ACKs with same number trigger a retransmission without a timeout (fast retransmission)

TCP SEGMENT FORMAT

- For information only, you will not be tested on the format of the header

Source port #										Destination port #									
Sequence Number																			
Acknowledgement Number																			
Hdr Len	Not used	CWR	ECE	URG	ACK	PSH	RST	SYN	FIN	Receiver Window									
Checksum										Urgent Data Pointer									
Options (optional, variable length)																			
Application payload (optional, variable length)																			

TCP CONNECTION ESTABLISHMENT

- TCP uses “three-way handshake” to establish connection
- Client sends initial SYN message
 - Initial sequence number for client→server is specified
- Server responds with SYN/ACK message
 - Client→server sequence number is confirmed in ACK
 - Server→client initial sequence number is specified
- Client sends an ACK message
 - Server→client sequence number is confirmed in ACK

TCP CONNECTION TERMINATION

- Side that wants to terminate sends FIN message
 - Other side responds with ACK
- Other side will also send a FIN message
 - It may not send it immediately, since it may have more data to send
 - Also responded with an ACK
- Alternative: connection abortion (RST message)
 - Usually used if other side misbehaves or if disconnection or too many timeouts are detected